

1.9 Deployment Architecture Documentation (*General Section*)

1. Server Placement

1.1 Production (Prod):

- The production environment is deployed to handle live traffic. It includes the application servers, databases, and all critical services running in a **highly available** configuration.
- The **K3S Cluster** will be deployed to the production environment, managing the microservices such as `Token-info-consumer`, `qpsdc-web`, `qshowcase`, `tkis-web`, and `QPA`.
- Services are deployed on a **private network** to prevent direct public access, with secure access points configured for internal services.

1.2 User Acceptance Testing (UAT):

- The UAT environment is a replica of the production setup, designed for testing new features and functionality.
- The **UAT server** will mirror the production environment in terms of **server configuration, service deployment, and database**. However, it may have limited data and fewer resources to simulate a production-like load.
- A **dedicated UAT database** will be maintained, separate from production, for testing purposes.

1.3 Disaster Recovery (DR):

- The DR environment is implemented to ensure business continuity in the event of system failures.
- This environment will be placed in a geographically **different data center** or cloud region to provide redundancy.
- The DR environment will replicate **core services** such as the **K3S Cluster, Token-info-consumer, MySQL database**, and other critical components.
- DR will be synchronized regularly with the production environment using **database replication** (covered below) and system snapshots for disaster recovery testing.

2. Load Balancing

2.1 Nginx Load Balancer (LB):

- **Nginx** will act as the primary **load balancer** for managing incoming traffic.
- Traffic will be distributed across multiple instances of the **Application Server**, ensuring efficient resource utilization and high availability.

- The load balancer will handle:
 - **Round-robin distribution** of HTTP requests to backend servers.
 - **SSL termination** for secure communication.
 - **Health checks** to monitor the availability of application servers and ensure automatic failover in case of server failure.
- The **Kong Ingress Controller** will handle API traffic routing, directing requests to the appropriate microservices based on the defined routes (e.g., /qshowcase, /QPA, /token-info-consumer).

2.2 Scaling Mechanism:

- The load balancing will automatically scale depending on traffic load, ensuring the system handles both **low and high traffic** efficiently.
- K3S provides **auto-scaling** based on the container resource usage and load, optimizing application performance.

3. Database Replication

3.1 MySQL Database Replication:

- To ensure high availability and **data consistency**, **MySQL** databases will be set up in a **master-slave replication** configuration.
- The **master database** will handle **read/write operations**, while **slave databases** will handle **read-only queries** to reduce the load on the master.
- **Replication** will be configured to ensure that data from the master is replicated in near-real-time to the slave database to avoid data loss in case of server failure.
- The replication setup will include automated failover mechanisms, so if the master database becomes unavailable, the system will promote one of the slaves to become the new master.
- The **DR environment** will have a replicated **MySQL instance** synchronized from the production master to ensure data availability during disaster recovery events.

3.2 Backup Strategy:

- **Automated backups** will be scheduled for both the master and slave databases at regular intervals to ensure recoverability of data in case of failure.
- **Point-in-time recovery** will be enabled for precise restoration of the database to a specific state.

4. Network Security

4.1 Firewall Configuration:

- All **application servers** and **database servers** will be placed behind **firewalls** to prevent unauthorized access.
- **Public-facing services** like the **Nginx load balancer** will have **restricted IP access**, allowing only trusted sources to reach the internal application and services.
- A **Virtual Private Cloud (VPC)** will be set up to ensure secure communication between services within the cloud environment.

4.2 SSL/TLS Encryption:

- All traffic between clients and servers will be encrypted using **SSL/TLS certificates**.
- The **Nginx load balancer** will handle SSL termination, ensuring secure communication from the client to the server.

4.3 Service-to-Service Communication:

- Communication between microservices like `qshowcase`, `tkis-web`, `Keycloak`, and **Token-info-consumer** will be secured using **mutual TLS** or **JWT-based authentication** to prevent unauthorized access.
- **API Gateway (Kong Ingress)** will enforce **authentication and authorization policies** for service-to-service communication.

4.4 Access Control & Authentication:

- **Keycloak** will manage the **authentication and authorization** of users accessing the system, ensuring proper user roles and access levels.
- Each service that requires access to the **database** or **other services** will be authenticated through **JWT tokens** issued by Keycloak.

5. Monitoring Components

5.1 System Monitoring:

- Tools like **Prometheus** and **Grafana** will be used for continuous **monitoring** and **visualization** of the system's performance.
- **Prometheus** will collect metrics such as:
 - CPU usage, memory usage, disk space.
 - Network latency and throughput.
 - API request/response times.

- **Grafana** will visualize these metrics and provide real-time **dashboards** to the operations team for monitoring system health.
- **Alerting** will be configured to notify the team in case of system failures or threshold breaches (e.g., high CPU usage, service downtime).

5.2 Log Monitoring:

- **Centralized logging** will be implemented using **ELK stack (Elasticsearch, Logstash, Kibana)** or **Fluentd** to aggregate logs from all services.
- Logs from the **application servers, microservices, and database** will be sent to the logging system for troubleshooting and auditing purposes.
- Alerts will be set up to notify the team in case of errors, failures, or anomalies in the logs.

5.3 Security Monitoring:

- **Security Information and Event Management (SIEM)** tools will be used to collect and analyze security-related data.

4.03 Diagram for Infrastructure Components (*Infra Preparation*)

Here are detailed solution diagram covering the infrastructure components for **Data Center (DC)**, **Disaster Recovery (DR)**, and **Test & Development (T&D)** environments. This diagram will showcase the entire infrastructure landscape, ensuring that all components are well-represented and align with the project's technical and business requirements.

1. Data Center (DC) Infrastructure Components

The **Data Center (DC)** environment represents the production infrastructure for the application, hosting the core services that support live traffic. This environment is designed for **high availability** and **scalability**.

A. Nginx Load Balancer (LB):

- Deployed in a highly available configuration to distribute traffic across multiple application servers, ensuring high availability and preventing service downtime.
- **Public-facing** for incoming traffic, configured with SSL/TLS for secure communication.

B. Application Servers (K3S Cluster):

- The core application components, including **microservices** like `Token-info-consumer`, `qpsdc-web`, `qshowcase`, `tkis-web`, and `QPA`, run on **K3S (Kubernetes)** for container orchestration and scaling.
- Microservices communicate via **Kong Ingress** for routing and API management.

C. Database (MySQL):

- The **MySQL database** is configured with **master-slave replication** to support high availability and load balancing between the read and write operations.
- The **master database** handles writes, while **slave databases** manage reads to optimize performance.

D. Authentication (Keycloak):

- **Keycloak** provides centralized authentication and authorization, integrated into the microservices architecture for managing user access and roles.

E. Message Queue (RabbitMQ):

- **RabbitMQ** serves as the messaging system for asynchronous communication between microservices, ensuring scalability and reliability.

F. **Cache (Redis):**

- **Redis** is used for caching frequently accessed data, improving performance and reducing the load on the databases.

1.B Network Setup:

- A. The **DC environment** will be deployed within a **Virtual Private Cloud (VPC)** with isolated subnets to ensure network security.
- B. Secure communication between services will be facilitated through **private IPs**, with **public IPs** only for necessary interfaces like Nginx load balancer and API Gateway.

2. Disaster Recovery (DR) Infrastructure Components

The **Disaster Recovery (DR)** environment is designed to ensure business continuity and quick recovery in case of system failure or data center unavailability. It is geographically isolated from the DC environment and synchronized regularly to maintain redundancy.

A. **Nginx Load Balancer (LB) - DR:**

- The DR environment will have a **secondary Nginx load balancer** to handle traffic routing during disaster recovery scenarios. It will operate as a **failover** if the primary data center fails.

B. **K3S Cluster - DR:**

- A mirrored **K3S cluster** will be deployed in the DR environment, replicating the microservices deployed in the DC.
- This ensures that, in case of a failure, the DR environment can handle live traffic with minimal downtime.

C. **Database Replication (MySQL):**

- The **MySQL database** in the DR environment will be configured to replicate from the master database in the DC environment.
- In the event of a disaster, the DR database will take over with the **most recent data** to ensure minimal data loss.

D. **Keycloak - DR:**

- A **backup Keycloak instance** will be deployed in the DR environment, synchronized with the production Keycloak for **authentication failover**.

E. **RabbitMQ - DR:**

- A **redundant RabbitMQ instance** will be set up in DR, ensuring that message queues are available even during a disaster scenario.

F. **Redis - DR:**

- Redis in the DR environment will provide caching redundancy, maintaining session state and critical data during failover.

2.B Network Setup:

- The DR environment will be set up in a different **cloud region** or **data center**, geographically separated from the DC to avoid simultaneous failures.
- Secure communication will be ensured between the DC and DR, with **VPNs** or **private direct connections** ensuring encrypted data transfer.

3. Test & Development (T&D) Infrastructure Components

The **Test & Development (T&D)** environment is designed for development, testing, and staging purposes. It is a replica of the DC environment but with fewer resources to allow developers and testers to validate new features without impacting production.

A. **Nginx Load Balancer (LB) - T&D:**

- The **Nginx load balancer** in the T&D environment will distribute traffic across the application servers used for development and testing.
- This setup ensures that new features and bug fixes are tested in an environment similar to production.

B. **K3S Cluster - T&D:**

- The **K3S Cluster** in T&D will mirror the production setup with similar configurations for microservices deployment, but with limited resources.
- It allows developers to test new changes with realistic load and scalability in mind.

C. **Database (MySQL) - T&D:**

- A **separate MySQL instance** will be deployed for T&D to prevent accidental data manipulation in production.
- **Data sanitization** and limited datasets will be used in this environment to comply with privacy and security regulations.

D. **Authentication (Keycloak) - T&D:**

- **Keycloak** will be set up in the T&D environment to handle authentication for users accessing the test and development services.

E. **Message Queue (RabbitMQ) - T&D:**

- A **RabbitMQ instance** will be deployed for message queueing during the development and testing process to simulate production-like conditions.

F. **Cache (Redis) - T&D:**

- **Redis** will be used in T&D for caching and session management during testing and performance evaluation.

3.B Network Setup:

- The **T&D environment** will be deployed in the same **cloud provider** but in a **separate VPC** or **subnet** to isolate it from production systems.
- The T&D network will have **restricted access** from production systems to ensure that no accidental changes affect live services.

4.04 Disaster Recovery (DR) Replication and Switch-Over Methods (*Infra Preparation*)

In this section, we outline our approach for **Disaster Recovery (DR)** replication and switch-over, which has been designed to minimize changes and ensure quick recovery with minimal downtime. This methodology includes the following key components:

1. Automated Data Replication

A. Database Replication:

- **MySQL Database:** We implement **master-slave replication** for the MySQL database to ensure that all data from the primary production database is automatically replicated to the disaster recovery (DR) database.
- **Replication Mode:** The **replication will be asynchronous**, allowing for efficient, low-latency replication of data to the DR environment.
- **Data Integrity:** Checks will be regularly performed to ensure the **consistency** between the master and the slave database, preventing data loss.
- **Automatic Failover:** In the event of failure, the replication setup ensures that the DR database can be promoted to the **master role**, allowing for continuous service availability with minimal disruption.

B. Real-Time Synchronization:

- To ensure that no significant data lag occurs during replication, **real-time synchronization** will be enabled.
- Continuous synchronization minimizes the risk of data loss during failover events.

C. Transactional Integrity:

- **Transactional logs** will be used to track all database changes to ensure that, in case of a failure, the DR database can recover to the exact state of the primary database at the time of failure.

2. Efficient Failover Mechanisms

A. Automated Failover Process:

- In case of a **disaster** or **service failure** in the primary environment, the failover process is fully **automated** to minimize manual intervention and reduce downtime.
- **K3S Cluster Failover:** The Kubernetes cluster running in DR will automatically detect any issues with the primary environment and switch the traffic to the DR environment.

- **Database Failover:** The **DR MySQL** instance will be automatically promoted to the master role, ensuring continued access to the most up-to-date data.

B. Application and Service Failover:

- The **Kong Ingress Controller** in the DR environment will automatically reroute requests to the available services without requiring manual configuration.
- **Nginx Load Balancer** will also support failover, detecting the unavailability of the primary servers and directing traffic to the backup instances in the DR environment.

C. Minimal Service Downtime:

- Failover processes are designed to be as **transparent** as possible to end-users, ensuring minimal service downtime. The goal is for failover events to result in **seconds** of downtime, if any.

3. Quick Recovery Procedures

A. Failover Testing:

- **Automated failover testing** will be conducted periodically to ensure that the failover process functions as expected during actual disaster events. This will include testing **data integrity**, **service availability**, and **application functionality** after failover.
- Failover testing will simulate failures at multiple points, including database, application servers, and load balancers.

B. Recovery Time Objective (RTO):

- The recovery time for each service will be carefully measured, and an **RTO of under 15 minutes** will be targeted for critical services, ensuring that business continuity is maintained with minimal disruption.

C. Backup and Restore Procedures:

- We will maintain **backup copies** of all critical data and configurations in the **DR environment**, ensuring that these can be quickly restored in the event of a failure.
- **Database snapshots** and **full backups** will be taken periodically to ensure that recovery can be performed at any time.

D. Rolling Recovery Strategy:

- In case of large-scale failure or infrastructure issues, the system will utilize a **rolling recovery** strategy where individual components (like application servers or

databases) are restored or rebuilt gradually, minimizing the impact on service availability.

4. Minimal Configuration Changes During Switch-Over

A. Consistent Configuration Across DC and DR:

- The **configuration management system** (e.g., **Helm charts** for Kubernetes or **Docker Compose** files) will ensure that the configuration between the **primary environment (DC)** and **disaster recovery (DR)** environment remains **identical**.
- This ensures that there are **minimal configuration changes** required during the switch-over to the DR environment, making the transition smooth and reducing the chances of misconfiguration.

B. Automated Configuration Sync:

- Configuration files for **Nginx**, **Kong**, and **microservices** will be synchronized between the **DC and DR environments**. The use of **Infrastructure-as-Code** (e.g., Terraform, Ansible) ensures that the environments are consistently deployed with the same configurations.

C. Load Balancer Configuration:

- The **Nginx load balancer** configuration will be automatically updated during the failover to ensure that traffic is directed to the correct backend services in the DR environment without needing manual intervention.

D. Service Discovery:

- Using **Kubernetes DNS** or **Consul** for service discovery ensures that microservices can find each other even in the event of a switch-over. No manual reconfiguration of service endpoints will be needed.

5. Simplified DR Testing Process

A. Non-Disruptive DR Testing:

- The **DR testing process** will be designed in a way that it does not disrupt the production environment. The testing will be conducted on a scheduled basis and will simulate failure scenarios without affecting live traffic.
- **Shadow Failover Testing:** We will perform “shadow” failover tests where the DR environment is activated in parallel with the live environment, and traffic is monitored to ensure the DR environment can handle it effectively.

B. Automated Testing Tools:

- We will leverage automated tools for testing the **failover procedures, replication health, and data consistency** between the primary and DR environments.
- These tests will simulate various failure scenarios, including network outages, database server crashes, and application server failures.

C. Testing Metrics and Reports:

- Detailed **logs and metrics** will be collected during DR tests, and a report will be generated for analysis to identify any areas of improvement.
- The results of these tests will help fine-tune the failover procedures to minimize downtime further and ensure the effectiveness of the DR strategy.

D. Continuous Improvement:

- Based on the results of DR tests, continuous improvements will be made to the **replication, failover, and recovery** mechanisms, ensuring the process is always optimized for the fastest recovery with the least impact.

4.06 High Availability Implementation for Data Center (*Infra Preparation*)

To ensure **high availability** (HA) within the Data Center (DC), we will implement various methods designed to minimize downtime, ensure service continuity, and enable quick recovery from any failure. The following methods will be used:

A. Load Balancing

- **Nginx Load Balancer (LB):**
 - **Nginx LB** will be deployed as the primary load balancer in the DC to manage incoming traffic and distribute it across multiple **application servers** and **Kubernetes pods**.
 - The load balancing mechanism will ensure that:
 - Traffic is efficiently distributed using **round-robin** or **least-connections** algorithms.
 - **SSL termination** is handled at the load balancer to secure communication.
 - **Auto-scaling** of application servers is supported, ensuring optimal resource usage during peak and off-peak times.
- **Kubernetes Ingress Controller:**
 - The **Kubernetes Ingress Controller** (Kong or Nginx Ingress) will manage API traffic routing to various microservices within the Kubernetes cluster.
 - **Path-based routing** will be used to direct traffic to specific services such as /qshowcase, /QPA, /token-info-consumer, etc., ensuring smooth traffic management within the cluster.
- **Session Affinity:**
 - **Sticky sessions** (session affinity) will be configured on the load balancer to ensure that a user's session remains connected to the same application server, maintaining state consistency.

B. Clustering

- **Kubernetes Cluster:**
 - The DC infrastructure will utilize a **Kubernetes (K8s) Cluster** to manage application services. The Kubernetes infrastructure will provide automatic

horizontal scaling of pods based on load, ensuring efficient use of resources and high availability.

- The cluster will be deployed across **multiple nodes** within the DC to ensure **resource redundancy** and prevent any single point of failure.
 - **StatefulSets** will be used for services like databases or Redis to ensure the persistence of state across pod restarts.
- **Microservices Clustering:**
 - Microservices like **qpsdc-web**, **qshowcase**, **QPA**, and **Token-info-consumer** will be deployed in **Kubernetes pods** with **replication controllers** to ensure multiple instances of each service are running across different nodes.
 - Kubernetes will automatically handle **pod scaling** based on traffic demands, ensuring the services remain highly available.

C. Automatic Failover

- **Kubernetes Auto-Healing and Pod Rescheduling:**
 - Kubernetes will automatically detect **failed pods** and **reschedule them** on healthy nodes to maintain service availability without human intervention.
 - If a **node fails**, Kubernetes will redistribute the affected pods across the remaining available nodes, ensuring the system continues to function without disruption.
- **Database Failover (MySQL):**
 - **MySQL replication** will be configured with **automatic failover** to ensure minimal downtime in case of database failures.
 - The **DR environment** will act as a backup for database failover, and **Kubernetes StatefulSets** will be used to ensure that the failover process does not cause data loss or service interruption.
- **Message Queue Failover (RabbitMQ):**
 - **RabbitMQ** will be deployed in a **clustered mode** to ensure message queue availability. If one RabbitMQ instance fails, the system will failover to another instance within the cluster, ensuring that messages are not lost and that services continue to operate.

D. Resource Redundancy

- **Kubernetes Node Redundancy:**
 - The **Kubernetes cluster** will be deployed across multiple physical or virtual nodes to ensure **node redundancy**. If one node fails, Kubernetes will reschedule the affected pods to available nodes without affecting the application's functionality.
 - **Persistent Storage Redundancy:**
 - Persistent storage (e.g., **MySQL databases**, **Redis** caches) will be replicated across multiple nodes or data centers to ensure data availability during hardware failures.
 - We will use **distributed storage** solutions like **Ceph** or **GlusterFS** for storage redundancy within Kubernetes.
- **Load Balancer Redundancy:**
 - **Nginx load balancers** will be deployed in a **highly available configuration** to ensure that, if one load balancer fails, traffic is automatically redirected to another load balancer without impacting user experience.
- **Keycloak Redundancy:**
 - **Keycloak** will be deployed in a **clustered mode** within Kubernetes to provide authentication and authorization services even if one Keycloak instance fails. Redundant replicas will be deployed to ensure consistent access control.

E. Service Health Monitoring

- **Prometheus and Grafana:**
 - **Prometheus** will continuously monitor the health of all **Kubernetes nodes, pods, and services** within the infrastructure, collecting metrics on:
 - **CPU usage, memory usage, disk space.**
 - **Response time and latency** of API services.
 - **Error rates and service uptime.**
 - **Grafana** will visualize these metrics in **real-time dashboards**, providing an overview of the cluster's health, enabling the Infra team to proactively detect and respond to issues.
- **ELK Stack (Elasticsearch, Logstash, Kibana):**

- The **ELK stack** will be used for centralized logging, ensuring that logs from all services (e.g., application, database, and infrastructure logs) are collected, indexed, and available for analysis.
- Logs will be automatically sent to **Elasticsearch** for fast querying and stored securely for auditing purposes.
- **Kibana** will provide real-time access to logs and enable Infra teams to quickly identify and address any potential issues affecting service health.
- **Loki & OpenTelemetry:**
 - **Loki** will be used alongside Prometheus for **log aggregation**, allowing Infra teams to correlate logs with metrics.
 - **OpenTelemetry** will be used to collect and report telemetry data such as **distributed tracing** for performance analysis, helping to identify bottlenecks or points of failure.
- **Librenms:**
 - **Librenms** will be integrated into the infrastructure for **network monitoring**. It will monitor network devices, servers, and services, sending alerts when any anomalies are detected, such as high network latency or device failures.

F. Fault Tolerance Mechanisms

- **Pod-level Fault Tolerance (Kubernetes):**
 - Kubernetes will provide **pod-level fault tolerance** by ensuring that critical services like **application servers**, **databases**, and **caches** are replicated across multiple nodes.
 - In case of node failure, Kubernetes will automatically reschedule pods to available nodes with minimal downtime.
- **Application-Level Fault Tolerance:**
 - **Circuit Breaker Patterns** will be used within application code to handle failures gracefully, preventing cascading failures and ensuring that the system remains functional even if certain services are temporarily unavailable.
 - **Retries and Timeouts** will be configured for service calls, ensuring that transient failures are automatically retried.
- **Service-Level Agreements (SLAs):**
 - We will define **SLAs** for all critical services to ensure that the application can meet availability and performance requirements even during periods of failure.

- Monitoring systems (Prometheus, Grafana) will track SLA adherence and trigger alerts if an SLA is breached.

4.07 Data Backup and Recovery Methods (*Infra Preparation*)

In this section, we outline the comprehensive **data backup** and **recovery** methods that will be implemented to ensure the **integrity**, **availability**, and **recoverability** of critical data in the event of a failure or disaster. These methods are designed to minimize downtime and ensure that we can restore the data to its most recent, consistent state.

A. Regular Automated Backups

- **Automated Backup Scheduling:**
 - **Automated backup schedules** will be set up to ensure that **critical data** (including databases, application files, and configurations) is backed up regularly without manual intervention.
 - Backups will occur at **predefined intervals** (e.g., daily, weekly) depending on the criticality and frequency of data changes.
- **Backup Tools:**
 - Tools like **Kubernetes Volumes Snapshots**, **MySQL Dump**, and **Cloud Backup Solutions** (e.g., AWS S3, GCP Cloud Storage) will be utilized for creating **consistent backups**.
 - These tools will be configured to run in the background, ensuring that backups are performed without interfering with regular operations.
- **Automated Backup Monitoring:**
 - **Monitoring systems** (e.g., Prometheus and Grafana) will be in place to ensure that backups are being completed successfully.
 - Alerts will be set up to notify the team if any backup fails, ensuring timely intervention.

B. Incremental and Full Backup Options

- **Full Backups:**
 - **Full backups** will be performed on a scheduled basis (e.g., weekly or monthly) to capture the entire data set, including all databases, files, and configurations.

- Full backups will serve as a reliable baseline from which other backup types (incremental and differential) are derived.
- **Incremental Backups:**
 - **Incremental backups** will be taken on a **daily basis** to back up only the **changes** made since the last backup.
 - This approach reduces storage space and backup time by focusing on new or modified data.
 - Incremental backups will be used to restore data to the most recent state without needing to restore the entire system from the full backup.
- **Differential Backups (Optional):**
 - **Differential backups** will capture the changes since the last full backup. They will be performed **less frequently** (e.g., every few days) to balance the backup size and speed.

C. Point-in-Time Recovery (PITR)

- **Database PITR:**
 - **Point-in-time recovery** will be implemented for databases such as **MySQL** using **binary logging**.
 - This enables the system to recover to any specific point in time by replaying transactions up to the desired moment.
 - PITR ensures minimal data loss in the event of system failures, allowing for recovery to the exact point in time before the failure occurred.
- **File and Configuration PITR:**
 - **File-based recovery** will be integrated to support point-in-time restoration for critical system files and configuration data.
 - This ensures that the entire environment, including databases and configurations, can be restored to a consistent state.

D. Backup Verification Processes

- **Automated Verification:**
 - Backups will undergo **automated verification** after completion to ensure the integrity and completeness of the backup data.
 - Verification will include checks for **data corruption**, **missing files**, and **incomplete backups**.

- **Test Restores:**
 - **Test restores** will be performed regularly to validate the integrity of the backup files. This process will ensure that backups can be restored successfully to both development and production environments.
 - Automated test restores will be scheduled on a **quarterly basis** to ensure that recovery processes are efficient and effective.
- **Alerting and Monitoring:**
 - Alerts will be configured to notify the team of any failures or discrepancies during the verification process. Any failed verification attempts will be immediately addressed by the system administrators.

E. Data Retention Policies

- **Retention Period:**
 - A **data retention policy** will be established to define how long different types of backup data (full, incremental, differential) will be retained.
 - For example, full backups will be retained for **1 year**, incremental backups will be retained for **6 months**, and older backups will be archived or deleted after this period.
- **Compliance with Regulations:**
 - Data retention policies will comply with relevant **regulations** and **industry standards** (e.g., GDPR, HIPAA) regarding the storage and retention of data.
 - Sensitive data will be stored securely and archived in a way that ensures compliance with data privacy laws.
- **Backup Archiving:**
 - Older backups will be moved to **long-term storage** solutions such as **AWS Glacier** or **GCP Coldline** for **cost-effective archiving**.
 - Archived data will be kept in a read-only state for compliance and recovery purposes.

F. Recovery Testing Procedures

- **Routine Recovery Drills:**
 - **Disaster recovery drills** will be conducted on a regular basis to ensure the team is prepared for any data loss scenario.

- These drills will involve simulating various failure scenarios and testing the restoration of backups to different environments, such as staging or production.
- Recovery time objectives (RTO) and recovery point objectives (RPO) will be measured during these drills to ensure compliance with SLAs.
- **Data Restoration Validation:**
 - After performing the recovery procedure, the restored data will be **validated** to ensure its integrity and consistency.
 - This will include verifying that databases, files, and configurations have been restored to their proper states and can function without errors.
- **Continuous Improvement:**
 - Post-recovery testing reviews will be conducted to identify any issues or bottlenecks in the backup and recovery process.
 - Any gaps or inefficiencies discovered during testing will be addressed to improve future recovery efforts.

5.04 OS and Platform Readiness Guide (*Virtualization & Operating System*)

A **detailed OS and platform readiness guide** that has been thoroughly tested and optimized for our application. This guide ensures that all necessary configurations, dependencies, and optimization settings are in place to guarantee **optimal performance** on the specified platforms.

A. Supported Operating Systems and Platforms

- **Operating Systems:**

- Our application is tested and optimized for the following operating systems:

- **Windows Server** (version?)
 - **Linux** (Ubuntu 20.04 LTS, CentOS 8, RHEL 8, or later)
 - **macOS** (for development and testing environments)

- **Virtualization Platforms:**

- The application is fully compatible with **VMware, Red Hat KVM, Oracle OVM**, and other major virtualization technologies.
 - Detailed instructions for configuring the application on virtualized platforms are included.

B. Installation and Setup Instructions

- **Pre-installation Requirements:**

- **System Requirements:** Minimum and recommended hardware specifications (CPU, RAM, disk space) for optimal performance.
 - **Network Configuration:** Ensure network interfaces are configured properly, including firewalls, DNS settings, and IP address allocation.
 - **Security Settings:** Ensure that proper security configurations are in place, including firewall settings and permissions for user roles.

- **Dependencies:**

- List of **software dependencies** such as:
 - Database (MySQL, PostgreSQL, etc.)
 - Web servers (Apache, Nginx, etc.)
 - Messaging systems (RabbitMQ, Kafka, etc.)
 - Any other required services or frameworks.

- Instructions for installing and configuring each dependency to ensure seamless integration with the application.
- **Installation Process:**
 - **Step-by-step installation guide** for setting up the OS and application on supported platforms.
 - **Automated installation scripts** or tools (if applicable) to streamline the setup process.

C. Configuration Settings

- **OS Configuration:**
 - **File System Settings:** Recommended file systems (e.g., ext4, XFS) and configurations (e.g., mount options).
 - **Swap and Memory Management:** Tuning settings for virtual memory, swap partitions, and **hugepages** configuration.
 - **Networking:** Configuration of network interfaces, including tuning for **high network throughput** and **low-latency** requirements.
 - **Security:** OS-specific security settings (e.g., disabling unused services, configuring SELinux/AppArmor, etc.).
- **Application Configuration:**
 - **Configuration Files:** Detailed breakdown of **configuration file parameters** for the application, such as **log paths**, **port numbers**, and **environment variables**.
 - **Performance Tuning:** OS-level and application-specific settings, including tuning of **max open files**, **worker threads**, and **heap sizes**.
- **Virtualization-Specific Configuration:**
 - Recommended settings for running the application in virtualized environments, including **vCPU**, **memory overcommitment**, and **disk I/O** optimizations.

D. Optimization Settings

- **Memory Optimization:**
 - Recommendations for optimizing **RAM usage** and **caching** to enhance application performance.
 - Configuring **NUMA** (Non-Uniform Memory Access) for multi-processor systems to improve memory access speed.

- **CPU Resource Management:**
 - Tuning CPU scheduling parameters for both **physical** and **virtual platforms**, including **CPU pinning** and **affinity settings** to ensure optimal distribution of resources.
- **I/O Optimization:**
 - File system tuning (e.g., **XFS** vs **ext4**), **disk caching** strategies, and **I/O scheduler** selection for optimal disk access performance.
 - Network optimizations, such as **TCP offloading**, **Jumbo frames**, and **network buffer tuning**.
- **Thread and Process Management:**
 - Recommendations for thread pool configurations, process scheduling, and managing high-concurrency workloads to reduce contention and improve throughput.

E. Testing and Validation

- **Performance Testing:**
 - Guidance on **benchmarking** the application's performance before and after configuration changes to verify improvements.
 - Recommended testing tools and frameworks to ensure that the platform is optimized for the application's specific requirements.
- **Compatibility Testing:**
 - **Integration testing** with all supported OS versions, ensuring the application runs smoothly in the specified environments without errors.
- **Failover and Recovery Testing:**
 - Steps to test the **resilience** of the platform, ensuring the application can handle failover scenarios and recover gracefully.

F. Troubleshooting and Support

- **Common Issues:**
 - A troubleshooting guide covering common configuration problems and their solutions, such as network or dependency issues.

- **Support Resources:**

- Links to official documentation, community forums, and support contact details for additional assistance if needed.

5.05 Required Changes for Users, Processes, and Firewall Configurations (*Virtualization & Operating System*)

We will disclose all the necessary changes related to **users**, **processes**, and **firewall configurations** required during the operation of the application. This information will be provided ahead of the **platform readiness** phase to allow for adequate planning, coordination, and implementation. These changes are essential to ensure that the application operates securely and efficiently in your environment.

1. User Configuration Changes

A. User Roles and Permissions:

- Detailed definition of **user roles** required for the application, including:
 - **Administrator:** Full access to all application functionalities and configurations.
 - **Power User:** Access to most features, but limited administrative capabilities.
 - **Standard User:** Access to specific application features based on business requirements.
- Ensure proper **role-based access control (RBAC)** is configured for securing access to sensitive data and operations.
- **User Authentication:**
 - If using an external authentication provider (e.g., **LDAP**, **Active Directory**, or **OAuth**), steps will be provided to integrate the application with these systems.
 - Configuration for **two-factor authentication (2FA)** or **multi-factor authentication (MFA)** if required.

B. User Creation and Management:

- Instructions for creating and managing user accounts, including **automated user provisioning** where applicable (e.g., using **SAML** or **OAuth**).
- Steps for managing **password policies** and **account lockout** policies to improve security.

C. Audit Logging:

- Enable **audit logs** to track user activity and ensure compliance with internal and external regulations.
- Logs should capture details of **user login attempts**, **access to sensitive data**, and any **administrative changes** made within the system.

2. Process Configuration Changes

A. Service and Process Management:

- **Process tuning** for both application and system services will be provided, including adjustments to:
 - **Service start-up** configurations to ensure that all essential services are launched during system boot.
 - **Process priority** settings for critical application components to ensure they receive sufficient CPU and memory resources.
- Process configuration files will be provided for both **Windows** and **Linux** systems, ensuring that the application runs optimally.

B. Background Jobs and Cron Jobs:

- A list of **scheduled tasks**, such as **cron jobs** for Linux or **Task Scheduler** for Windows, needed to manage routine application operations (e.g., backups, report generation, etc.).
- Configuration of background jobs to ensure **efficient resource utilization** without overloading the system.

C. Service Dependencies:

- Details on **process dependencies**, ensuring that all required services (e.g., databases, messaging queues) are up and running before the application processes start.
- Information on **service restart policies** to minimize downtime during failure scenarios.

D. Monitoring and Health Checks:

- A list of recommended **health checks** to monitor essential processes, such as the application server, database, and networking services.
- Tools for **automated process restarts** in case of failures or resource exhaustion.

3. Firewall Configuration Changes

A. Network Ports and Protocols:

- A comprehensive list of **network ports** and protocols required by the application, ensuring the firewall allows appropriate traffic:
 - **HTTP/HTTPS**: For web traffic.
 - **Database Ports**: For communication with the database (e.g., **MySQL port 3306**, **PostgreSQL port 5432**).
 - **Message Broker Ports**: For systems like **RabbitMQ**, **Kafka**, or similar message brokers.
 - **SSH**: For secure remote management (if applicable).
- Specific port ranges or IP addresses may be needed for cloud environments or multi-tier applications (e.g., **load balancer**, **application server**, **database server**).

B. Inbound and Outbound Traffic:

- Configurations for **inbound traffic** (e.g., allowing external connections to the web server, database access, etc.).
- **Outbound traffic** rules for services like **email servers** (SMTP), third-party APIs, or external integrations.

C. Virtual Private Network (VPN):

- VPN configurations for **secure remote access** to application components, especially for management purposes.
- Instructions for setting up a **site-to-site VPN** or **client VPN** to ensure secure communications between different sites or remote users.

D. Firewall Rules for Virtualized Environments:

- Specific firewall configurations for **virtualized platforms** (VMware, Red Hat KVM, Oracle OVM) to ensure that VMs communicate securely with each other and with external resources.
- Ensure that virtual networks or subnets are appropriately segmented and isolated, providing **secure communication** while avoiding unauthorized access.

4. Pre-Deployment Checklist

A. User Management:

- Ensure that all user accounts and roles are set up and properly configured.
- Verify that **RBAC** is implemented as per organizational security policies.

B. Service and Process Validation:

- Confirm that all necessary services and background jobs are configured correctly and have been validated for proper execution.

C. Firewall Configuration Verification:

- Verify that all required ports are opened and secured based on the application's network traffic needs.
- Perform tests to confirm **proper connectivity** between application components and external services.

D. Access Control and Security:

- Verify that **user authentication, audit logging, and password policies** are configured and functioning correctly.
- Ensure that **firewall settings** prevent unauthorized access while allowing necessary traffic.

6.04 Backup and Restoration Requirements (*Storage & Backup*)

We will provide a **comprehensive backup and restoration strategy** that aligns with the business owner's **backup frequency** and **retention policy**. This strategy will ensure that critical data, including **databases**, **configuration files**, and **application data**, are safely backed up and can be restored when needed. The backup system will be designed to maintain high availability, security, and efficiency.

1. Backup Frequency and Retention Policy

A. Backup Frequency:

- **Database Backups:** Regular **daily** backups will be performed for the databases to ensure that the latest data is protected. Depending on the business requirements, more frequent backups (e.g., **hourly**) can be implemented for high-priority databases.
- **Configuration Files:** **Weekly** backups of all critical configuration files will be taken to ensure that any changes can be rolled back in case of an issue.
- **Application Data:** Backups for application data will be conducted on a **daily** basis, with more frequent backups available upon request for critical data sets.

B. Retention Policy:

- **Database Backups:**
 - **Daily backups** will be retained for **7 days**.
 - **Weekly backups** will be retained for **1 month**.
 - **Monthly backups** will be retained for **1 year**.
- **Configuration Files:** Configuration file backups will be retained for a **minimum of 6 months** to allow for historical rollback of any configuration changes.
- **Application Data:**
 - Backups will be retained based on the importance and size of the data, with **daily backups** kept for **7 days**, **weekly backups** for **1 month**, and **monthly backups** for up to **1 year**.

2. Backup Process

- **Database Backups:**
 - We will use **full, incremental, or differential backup strategies** based on the size and criticality of the database:

- **Full backups** will be taken on a regular basis to capture the entire database.
 - **Incremental backups** will be taken to capture only changes made since the last backup, ensuring efficient use of storage space.
 - **Point-in-time recovery (PITR)** will be enabled, allowing for database restoration to a specific moment in time.
- **Backup Locations:** Backups will be stored securely in designated storage systems such as **on-premise servers**, **cloud storage** (e.g., AWS S3, GCP), and **disaster recovery** sites.

C. Configuration Files and Application Data Backup (Codelab):

- **Codelab Backup:** Configuration files and application data will be backed up in the **Codelab environment**, ensuring that the latest versions of application configurations and essential data are available for restoration.
- **Backup Mechanism:** Codelab will implement a versioned backup strategy, ensuring that both current and historical versions of configuration files and application data are available. Backups will be scheduled to ensure minimal impact on system performance.
- **Backup Storage:** Data in Codelab will be securely stored, with appropriate encryption both in transit and at rest. The backup will include:
 - **Configuration files:** Application, database, and system configurations that define the environment.
 - **Application data:** Files or databases required for the application to function properly (e.g., user data, logs, settings).

D. Automated Backup Procedures:

- We will implement automated scripts for backup processes, ensuring **consistent and reliable backups** of databases, configuration files, and application data.
- **Backup Monitoring:** Automated monitoring will be in place to ensure that backups complete successfully and alert the system administrator if any issues are encountered.

3. Restoration Process

- **Database Restoration:**
 - In the event of data loss or corruption, databases will be restored from the most recent **full backup** or **incremental backup**, ensuring minimal data loss.
 - Point-in-time recovery will be supported, enabling the restoration of databases to a specific time, down to the minute.

- **Configuration Files and Application Data Restoration (Codelab):**
 - **Codelab-based Recovery:** Configuration files and application data stored in Codelab can be easily restored to their latest state or a specific historical version.
 - **A version control system** will be implemented for configuration files to allow easy rollback to any previous version without downtime.
 - **Application Data Recovery:** Application data backups will be restored by accessing the relevant backup version, ensuring that the system can return to an operational state without data loss.
- **Disaster Recovery:**
 - In the event of a complete system failure, backup restoration will follow the **disaster recovery plan**, where the system is rebuilt using the latest backups.
 - **Failover and Redundancy:** Redundant systems will be in place to ensure **high availability** during the restoration process, allowing services to be restored with minimal downtime.

4. Security and Compliance

- **Encryption:**
 - All backups will be encrypted to ensure **data confidentiality**. Encryption will be applied both during transmission (using TLS) and while stored (using AES-256).
- **Access Control:**
 - Strict **access control policies** will be enforced for backup data to ensure that only authorized personnel have access to restore or manage backups.
 - Regular audits will be conducted to ensure compliance with security policies and to verify the integrity of backup data.
- **Compliance:**
 - The backup process will comply with relevant industry standards and regulations, including **GDPR**, **HIPAA**, and **SOX** (if applicable), ensuring data is handled and retained properly.

7.12 Database Account and User Management Approach (*Database*)

1. Access Control Policies

A. Internal Access Control:

- We enforce **strict access control policies** to ensure that only authorized personnel can access the database and sensitive application data.
- Access is granted based on the **principle of least privilege**, ensuring that users have only the necessary permissions to perform their duties.
- **Access auditing** and **logging** are enabled to track and monitor any attempts to access sensitive data, ensuring accountability for user actions.

B. Role-Based Access Control (RBAC):

- We use **role-based access control (RBAC)** to define and manage user roles and permissions within the database.
- Users are assigned specific roles that align with their job responsibilities, limiting their access to only the data and actions relevant to their work.
- Example roles include **Administrator**, **Read-Only User**, and **Data Analyst**, each with distinct permissions based on operational needs.
- Administrative roles are tightly controlled and require **multi-factor authentication (MFA)** to further secure access.

C. Separation of Duties:

- A **separation of duties** policy is applied to ensure that no single individual has the ability to both initiate and approve sensitive changes in the database, reducing the risk of fraudulent or unauthorized actions.

2. Secure Authentication

A. Password Policies:

- We implement **strong password policies** that mandate complexity and expiration to reduce the risk of password-based breaches. Passwords are regularly rotated and must meet minimum security standards, such as a combination of **uppercase**, **lowercase**, **numerals**, and **special characters**.

B. Multi-Factor Authentication (MFA):

- We require **multi-factor authentication (MFA)** for all privileged accounts, ensuring that access to the database requires both **something the user knows** (password) and **something the user has** (e.g., token or biometric verification).
- This significantly reduces the likelihood of unauthorized access due to compromised credentials.

C. Secure User Authentication:

- For third-party and remote access, we implement **VPN** and **SSL/TLS encryption** to ensure that data transmitted between the user and the database is secure.

3. Secure Data Access

A. Encryption of Data in Transit:

- All sensitive data transferred between the database and external systems or users is encrypted using **SSL/TLS** to prevent unauthorized interception.

B. Encryption of Data at Rest:

- Sensitive data, such as personally identifiable information (PII) or payment details, is encrypted at rest to ensure that it remains secure even in the event of physical access to storage media.
- Encryption keys are managed securely, with periodic rotations and strong access controls to prevent unauthorized access to encrypted data.

4. Monitoring and Auditing

A. User Activity Logging:

- Comprehensive **logging** of all database activity is implemented, tracking actions such as data access, changes to configurations, and administrative operations.
- These logs are securely stored and regularly reviewed to detect any anomalies or unauthorized access attempts.

B. Audit Trails:

- We maintain detailed **audit trails** of all user actions related to database access and management. These trails ensure accountability, providing a clear record of who accessed what data and when.
- **Alerting mechanisms** are in place to notify administrators of suspicious activities or policy violations in real-time.

5. Compliance and Security Improvements

A. Continuous Monitoring:

- We continuously review and improve our internal security practices to align with industry standards and best practices.
- Regular **security assessments**, including **penetration testing** and **vulnerability scanning**, are performed to ensure that our security measures are effective.

B. Ongoing PCI DSS Compliance:

- While our current system does not fully meet **PCI DSS** standards, we are actively working on aligning our database account and user management approach with these standards as part of our ongoing commitment to data security.
- We are conducting regular **PCI DSS gap assessments** to identify areas for improvement and ensure that we meet compliance requirements as we evolve our security strategy.